

Transcript Evaluator - Batch Report

Run time: 2026-05-30T19:47:06.629Z

Model: gemma-4-26b-a4b-it

Input file: input data/entire_interview.json

Batch size: 5

Total runtime: 927.52s

Status: Completed | Total: 22 | Success: 22 | Failed: 0 | Avg overall: 4.72

Case #2: Auto case 2

Rubric: behavioral

Resolved path: behavioral

Overall: 5/5

Case runtime: 29.22s

Question: Thanks. Let's go chronologically — walk me through your resume, starting from education or your first relevant role, whatever you think is the right entry point.

Response: I'll start with university briefly, then the jobs that matter for this role. I studied computer science at State University of New York at Buffalo — graduated 2017. My first job was at CardNest, a small payments startup in New York, for about two years. I was mostly doing CRUD APIs in Ruby, but I got exposure to PCI scope and why you don't log certain fields. In 2019 I moved to PayLoom, still mid-stage — that's where I switched to Go because the team was rewriting a monolith into services. I owned the merchant webhook delivery service: retries, signing, delivery guarantees. We scaled from a few thousand merchants to low six figures of daily events while I was there. I left in early 2022 when LedgerPay recruited me for reconciliation — they needed someone who had seen high-volume money movement and wasn't scared of batch windows. At LedgerPay I've been senior engineer for three years, tech lead on reconciliation for the last eighteen months. On paper my title is still senior, but I've been acting lead on the settlement migration and on-call rotation. Happy to go deeper on any of those stops.

Follow-up question: On PayLoom — you said low six figures of daily webhook events. What broke first when you scaled?

Follow-up response: Honestly, our Postgres write pattern on the delivery log table. We were inserting one row per attempt without partitioning, and hot indexes killed us around eighty thousand events a day. I proposed monthly partitions and moving retry scheduling to a Redis-backed queue so we weren't polling the database. We shipped partitions first — that bought headroom — then the queue refactor over two sprints. p99 delivery latency went from about four seconds to under eight hundred milliseconds at peak. That's the project I'm proudest of from PayLoom because it was the first time I owned a production scaling problem end to end.

Scorable

Yes - The candidate provides a detailed career history with specific technical ownership and measurable outcomes.

Dimensions

base (5/5): The candidate provided a clear, structured, and highly relevant chronological overview of their professional experience.

personal_contribution (5/5): The candidate clearly articulated their ownership of the webhook service and their specific role in proposing and implementing the architectural fix.

real_example (5/5): The candidate provided specific technical details regarding Postgres write patterns, indexing issues, and the transition to a Redis-backed queue.

outcome (5/5): The candidate provided a quantifiable and impressive result regarding p99 delivery latency reduction.

Follow-up

present=yes | probe=deepening | impact=positive_signal | The candidate provided significant technical depth and specific metrics in response to the probe, reinforcing their credibility.

Overall Reasoning

The candidate demonstrated strong technical ownership, provided specific and quantifiable outcomes, and successfully handled a deepening probe with high-quality technical detail.

Flags

None

Case #5: Auto case 5

Rubric: behavioral

Resolved path: behavioral

Overall: 4/5

Case runtime: 59.19s

Question: Good. Why Northline, and why this role specifically?

Response: Two reasons. First, your public write-up on the event-sourced ledger and how you handle out-of-order settlement events — that's exactly the class of consistency problems I've been working on, and I want to go deeper at your scale. Second, the job description called out on-call rotation and production ownership for platform engineers, not a handoff-to-ops model. I've run primary on-call for reconciliation for eighteen months and I think owning what you ship is how you build good systems. I'm not looking for a feature factory role; I want to be closer to the core transaction path.

Scorable

Yes - The candidate provides a substantive and highly specific response to the question.

Dimensions

base (5/5): The response is articulate, well-structured, and demonstrates deep research into the company.

personal_contribution (4/5): The candidate clearly identifies their specific expertise in consistency problems and their experience in on-call roles.

real_example (4/5): The candidate references specific technical domains like event-sourced ledgers and reconciliation processes.

outcome (2/5): The candidate describes their experience and tenure but does not explicitly state a specific result or impact achieved.

Follow-up

present=no | probe=none | impact=none | No follow-up probe was provided in the input.

Overall Reasoning

A very strong response that demonstrates high technical alignment and relevant experience, though it lacks specific outcome metrics.

Flags

None

Case #1: Auto case 1

Rubric: non_behavioral

Resolved path: non_behavioral

Overall: 5/5

Case runtime: 68.72s

Question: Let's start broad. Tell me about yourself — the version you'd give if we only had two minutes.

Response: Sure. I'm a backend engineer with about seven years in fintech and payments. Most recently at LedgerPay I owned the settlement reconciliation service — batch jobs, idempotency, the unglamorous stuff that has to be correct. What I care about is building systems that stay boring under load: clear contracts, good observability, and failure modes you can reason about in an incident. Before LedgerPay I was at a smaller startup where I wore more hats, which is where I picked up Go and got comfortable owning services end to end. I'm here because Northline's platform team sits on the critical path for fraud and ledger consistency, and that's the kind of problem I want my next few years to be about.

Scorable

Yes - The candidate provides a clear and professional summary of their background.

Dimensions

base (5/5): The response is well-structured, concise, and effectively connects past experience to the target role.

technical_relevance (5/5): The candidate uses specific, high-value technical concepts like idempotency and observability that are relevant to the role.

domain_alignment (5/5): The candidate demonstrates deep familiarity with fintech-specific challenges like settlement and ledger consistency.

Follow-up

present=no | probe=none | impact=none | None

Overall Reasoning

A highly effective professional pitch that demonstrates both technical authority and strong alignment with the company's domain.

Flags

None

Case #4: Auto case 4

Rubric: non_behavioral

Resolved path: non_behavioral

Overall: 5/5

Case runtime: 84.65s

Question: Got it. Any gap on the resume we should address? I don't see anything long between roles.

Response: No intentional gaps. There's a three-week overlap in 2022 when I finished PayLoom and started LedgerPay — I used that for relocation and to finish an online course on distributed tracing, which I've used heavily since.

Scorable

Yes - The candidate provides a direct and informative response to a specific question about their employment history.

Dimensions

base (5/5): The candidate provides a clear, proactive explanation that addresses the question and highlights professional development.

Follow-up

present=no | probe=none | impact=none | None

Overall Reasoning

The candidate provided a direct, transparent, and value-added response to a clarifying question.

Flags

None

Case #3: Auto case 3

Rubric: non_behavioral

Resolved path: non_behavioral

Overall: 5/5

Case runtime: 86.29s

Question: And LedgerPay — you list reconciliation and settlement. Are those the same system or separate?

Response: Separate services, shared data model. Reconciliation compares internal ledger entries to bank files and flags exceptions. Settlement is the outbound generation — ACH files, NACHA formatting, cutoffs. They talk through an event bus; settlement emits "file submitted" and reconciliation consumes it for next-day matching. I own reconciliation; I was tech lead on the settlement migration last year but day-to-day ownership is split with another lead now.

Scorable

Yes - The candidate provides a detailed technical explanation of the systems mentioned.

Dimensions

base (5/5): The candidate provides a clear, accurate, and technically detailed distinction between the two systems.

technical_depth (5/5): The candidate demonstrates deep domain knowledge by mentioning specific protocols, file formats, and architectural patterns.

Follow-up

present=no | probe=none | impact=none | No follow-up was provided in the input.

Overall Reasoning

The candidate provided a highly precise and technically sophisticated answer that clearly distinguished between the two systems.

Flags

None

Case #8: Auto case 8

Rubric: behavioral

Resolved path: behavioral

Overall: 5/5

Case runtime: 24.68s

Question: Different topic — describe a situation where you had to persuade a skeptical stakeholder.

Response: At LedgerPay last year, product wanted to defer investing in dead-letter queue visibility for our async payout pipeline to hit a marketing launch. Our VP of Engineering was skeptical it was worth two sprints. I put together a short doc: three production incidents in the prior quarter where we lost two-plus hours each because we couldn't see poison messages, an estimate that DLQ tooling would cut mean time to recovery by roughly half based on those incidents, and a phased plan to ship read-only inspection in one sprint and alerting in the second. I walked her through it in a thirty-minute slot and addressed the roadmap concern by showing the second sprint could run parallel to the launch freeze. She approved both sprints. We shipped the dashboard before Black Friday and it was used in a real incident that week — we cleared a stuck payout batch in under twenty minutes instead of the usual multi-hour log dive.

Scorable

Yes - The candidate provides a complete behavioral response with a clear situation, action, and measurable outcome.

Dimensions

base (5/5): The response is highly structured, concise, and directly addresses the prompt.

personal_contribution (5/5): The candidate clearly defines their specific actions, including data analysis, document creation, and strategic negotiation.

real_example (5/5): The candidate provides specific context, including the company, the technical problem, and the specific stakeholders involved.

outcome (5/5): The outcome is quantified with specific time-to-resolution improvements and a successful real-world application.

Follow-up

present=no | probe=none | impact=none | No follow-up probe was provided in the input.

Overall Reasoning

This is a top-tier behavioral response that demonstrates data-driven influence and clear business impact.

Flags

None

Case #9: Auto case 9

Rubric: behavioral

Resolved path: behavioral

Overall: 5/5

Case runtime: 26.62s

Question: Tell me about a time you had a conflict with a teammate.

Response: So, um, at LedgerPay in mid-2022 I disagreed with a senior engineer about test strategy for a new ledger adjustment service. He wanted integration tests only to move faster. I was uncomfortable because the service touched balance mutations. I asked for a one-on-one instead of debating in the team channel — I didn't want it to feel public. I walked through concrete failure modes unit tests would catch that integration tests had missed on a sister service. We landed on a hybrid: unit tests on the mutation layer, integration tests for the HTTP and messaging boundaries. The service launched with no adjustment-related incidents in the first two months.

Scorable

Yes - The candidate provides a complete and structured response to a behavioral question.

Dimensions

base (5/5): The response is professional, well-structured, and directly addresses the prompt.

personal_contribution (5/5): The candidate clearly outlines their specific actions: choosing a private setting and providing technical evidence.

real_example (5/5): The candidate provides specific context including company name, timeframe, and technical details.

outcome (5/5): The candidate describes both a technical compromise and a successful business result with measurable stability.

Follow-up

present=no | probe=none | impact=none | No follow-up probe was provided in the input.

Overall Reasoning

The candidate delivered a high-quality STAR response with clear technical reasoning and measurable impact.

Flags

None

Case #7: Auto case 7

Rubric: behavioral

Resolved path: behavioral

Overall: 5/5

Case runtime: 28.14s

Question: Makes sense. Let's shift to behavioral examples. Tell me about a time you led a project under a tight deadline.

Response: In Q3 2023 at LedgerPay we had to migrate our ACH settlement file generator to a new bank format before a regulatory cutoff in six weeks. I was tech lead for a team of four. The stakes were a potential processing halt if we missed the date. I mapped every field mapping and exception path in week one, split work into schema validation versus production cutover, and ran daily fifteen-minute standups focused only on blockers. I took the riskiest piece — the idempotent replay logic — myself. On day twelve we found the bank's sandbox dropped duplicate trace numbers; I escalated to their integration team and got a documented workaround. We went live on day thirty-eight, two days early, and the first production week had zero settlement exceptions tied to the migration.

Follow-up question: You mentioned idempotent replay. What would have broken if you hadn't gotten that right?

Follow-up response: We would have double-posted or skipped settlements on retry. In our system that surfaces as customer balance drift, which is a Sev-1. The replay layer keys on bank trace ID plus file sequence so a partial failure can resume without re-applying successful rows. We proved it with chaos tests that killed the worker mid-batch before we cut over.

Scorable

Yes - The candidate provides a complete behavioral response with a clear follow-up that adds significant technical depth.

Dimensions

base (5/5): The response is highly structured, addressing the situation, task, action, and result clearly.

personal_contribution (5/5): The candidate explicitly details their specific actions, such as mapping fields and taking on the riskiest technical task.

real_example (5/5): The response uses specific context, including dates, company names, and technical constraints.

outcome (5/5): The outcome is quantified with a specific timeline and a zero-error metric.

Follow-up

present=yes | probe=deepening | impact=positive_signal | The candidate added significant new technical detail regarding system impact and validation methods.

Overall Reasoning

An exceptional response that demonstrates leadership, technical ownership, and measurable success, further validated by a high-quality technical follow-up.

Flags

None

Case #10: Auto case 10

Rubric: behavioral

Resolved path: behavioral

Overall: 5/5

Case runtime: 28.77s

Question: Tell me about a time you failed or made a significant mistake.

Response: In 2023 I merged a change that tightened a validation rule on settlement files without realizing an older bank integration still sent a deprecated optional field. Production rejected about twelve hundred rows over a weekend batch window. I was primary on-call. I rolled back within twenty minutes, but the batch had already partially committed so we had to run a manual repair script Monday morning with finance watching. I owned the postmortem. Root cause was I didn't check integration-specific fixtures — we had tests for the new format only. After that I added a contract test suite per bank partner and made "partner matrix" a required checkbox on the PR template for anything touching file schemas. We haven't had a repeat, but I still think about it when someone says a change is "just validation."

Scorable

Yes - The candidate provides a complete behavioral response including context, specific actions, and long-term outcomes.

Dimensions

base (5/5): The response is highly structured, clear, and directly addresses the prompt.

personal_contribution (5/5): The candidate clearly identifies their specific roles in the rollback, the manual repair, and owning the postmortem.

real_example (5/5): The candidate provides specific technical details, such as the number of rows affected and the exact nature of the validation error.

outcome (5/5): The candidate describes both the immediate resolution and the systemic improvements made to prevent recurrence.

Follow-up

present=no | probe=none | impact=none | No follow-up probe was provided in the input.

Overall Reasoning

A top-tier behavioral response that demonstrates accountability, technical depth, and a proactive approach to systemic improvement.

Flags

None

Case #6: Auto case 6

Rubric: behavioral

Resolved path: behavioral

Overall: 5/5

Case runtime: 47.86s

Question: How have you used Go in production — not just syntax, but how you structure services?

Response: At LedgerPay almost everything I ship is Go. Typical layout: cmd/ for binaries, internal/ for domain packages, thin HTTP handlers that delegate to a service layer. I use context for cancellation everywhere, structured logging with trace IDs propagated from our ingress, and table-driven tests on pure logic. Concurrency-wise I prefer errgroup for bounded parallel work — like parsing bank file rows in workers with a fixed pool — rather than unbounded goroutines. I've also learned to be disciplined about not leaking goroutines on retries; we had an incident once from a forgotten timer channel that taught the team to always tie worker lifetimes to context cancel. I'm not a generics power user, but interfaces for storage boundaries — file store, ledger client — make testing with fakes straightforward.

Scorable

Yes - The candidate provides a substantive response with specific technical details and a real-world example.

Dimensions

base (5/5): The candidate demonstrates expert-level knowledge of Go production standards, directory structures, and concurrency patterns.

personal_contribution (5/5): The candidate clearly articulates their specific architectural decisions, testing strategies, and concurrency preferences.

real_example (5/5): The candidate provides a specific, technical example of a production incident involving a goroutine leak caused by a timer channel.

outcome (5/5): The candidate describes a successful resolution that resulted in a permanent change to team practices to prevent future leaks.

Follow-up

present=no | probe=none | impact=none | No follow-up was provided in the input.

Overall Reasoning

The candidate delivered a high-signal response that demonstrated technical depth, personal agency, and the ability to learn from mistakes.

Flags

None

Case #11: Auto case 11

Rubric: behavioral

Resolved path: behavioral

Overall: 5/5

Case runtime: 23.54s

Question: Tell me about a time you received tough feedback and what you did with it.

Response: Early in the settlement migration my director told me I was over-explaining in cross-functional updates and losing the room — product and finance stopped asking questions because my updates ran twenty minutes. I was defensive at first because I thought detail meant rigor. I asked for a concrete example and she pointed to a steering meeting recording. I watched it and she was right: I buried the decision. I switched to a format: decision needed, recommendation, risks, appendix doc for depth. I also started sending pre-reads twenty-four hours ahead. Next steering meeting we finished in twelve minutes and we got sign-off on the cutover plan that week. The feedback stung but it made me a better tech lead.

Scorable

Yes - The candidate provides a complete behavioral response with a specific situation, action, and result.

Dimensions

base (5/5): The response is highly structured, addressing the feedback, the reaction, the action, and the result clearly.

personal_contribution (5/5): The candidate took proactive steps to improve by seeking concrete examples and implementing a new communication framework.

real_example (5/5): The candidate uses a specific context regarding a settlement migration and steering meetings.

outcome (5/5): The candidate provides a clear, positive outcome involving improved meeting efficiency and project sign-off.

Follow-up

present=no | probe=none | impact=none | No follow-up probe was provided in the input.

Overall Reasoning

This is an excellent behavioral response that demonstrates self-awareness, coachability, and measurable professional growth.

Flags

None

Case #13: Auto case 13

Rubric: behavioral

Resolved path: behavioral

Overall: 5/5

Case runtime: 30.93s

Question: Tell me about a time you mentored someone or helped them grow.

Response: I paired with a junior engineer twice a week for a month. I gave them structured code review feedback with explanations rather than corrections. I also assigned them a small feature end-to-end so they'd have full ownership. By the end of the month they were submitting PRs that needed minimal review. They were promoted six months later.

Scorable

Yes - The candidate provides a complete behavioral response with specific actions and a clear outcome.

Dimensions

base (5/5): The response is direct, coherent, and fully addresses the prompt.

personal_contribution (5/5): The candidate explicitly details their specific actions, such as pairing, feedback style, and task delegation.

real_example (5/5): The candidate describes a specific instance involving a junior engineer and a defined timeframe.

outcome (5/5): The candidate provides both immediate performance improvements and a long-term career milestone.

Follow-up

present=no | probe=none | impact=none | No follow-up question was provided in the input.

Overall Reasoning

The candidate delivered a high-quality behavioral response with clear actions and tangible results.

Flags

None

Case #15: Auto case 15

Rubric: behavioral

Resolved path: behavioral

Overall: 5/5

Case runtime: 31.38s

Question: What is your greatest weakness?

Response: I can be slow to delegate when I'm under pressure on a critical path. It showed up last year during the settlement migration — I took on too many code reviews myself and became a bottleneck for the team for about a week. My manager flagged it in a one-on-one. Since then I've been explicit about pairing juniors on reviewable chunks earlier, and I use a simple rule: if I'm the only person who can approve more than three PRs in a day, I stop coding and redistribute. It's still a tendency when stakes feel high, but the last two projects hit review SLA without me being the gate.

Scorable

Yes - The candidate provides a complete, structured response to a behavioral-style question.

Dimensions

base (5/5): The candidate demonstrates high self-awareness by identifying a specific, non-cliché weakness.

personal_contribution (5/5): The candidate implemented a concrete, rule-based system to manage their tendency to bottleneck.

real_example (5/5): The candidate provided a specific context (settlement migration) and a clear consequence (one-week bottleneck).

outcome (5/5): The candidate cited a measurable improvement in meeting review SLAs in recent projects.

Follow-up

present=no | probe=none | impact=none | No follow-up probe was provided in the input.

Overall Reasoning

This is an excellent behavioral response that follows a clear situation-action-result pattern with high specificity.

Flags

None

Case #14: Auto case 14

Rubric: non_behavioral

Resolved path: behavioral

Overall: 5/5

Case runtime: 73.43s

Question: I'd like a bit more context on that one — where was this, and why were you mentoring them?

Response: That was at LedgerPay in 2023. She joined the reconciliation team straight from bootcamp. Her manager asked me because I had bandwidth after we shipped the partition work. The gap was system thinking — she could fix tickets but didn't see how reconciliation tied to finance's month-end close. I had her own the "exception reason codes" feature: small surface area but touched schema, API, and a finance report. I did office hours twice a week and required her to write the rollout plan. Promotion six months later was partly on that feature shipping without rollback.

Scorable

Yes - The candidate provides a highly detailed and specific response to a follow-up question regarding a past experience.

Dimensions

base (5/5): The response is structured, concise, and directly addresses the context requested.

personal_contribution (5/5): The candidate explicitly details their actions, such as conducting office hours and setting specific requirements for the mentee.

real_example (5/5): The candidate provides specific details including the company name, year, team, and the exact feature name.

outcome (5/5): The candidate provides a clear, measurable outcome: a promotion and a successful, rollback-free deployment.

Follow-up

present=yes | probe=deepening | impact=positive_signal | The candidate added significant new details that provided much-needed depth to the previous context.

Overall Reasoning

This is an excellent behavioral response that provides specific context, clear personal actions, and a tangible outcome.

Flags

None

Case #12: Auto case 12

Rubric: behavioral

Resolved path: behavioral

Overall: skip

Case runtime: 697.31s

Question: Tell me about a time you improved a process.

Response: At my last job we had a slow deployment process. I looked at it and figured out we could automate the manual steps with a script. I wrote the script myself, tested it, and rolled it out to the team. It ended up saving everyone a lot of time.

Follow-up question: Can you be more specific — what was slow, what did the script do, and how much time did you save?

Follow-up response: Fair push. The pain was our staging promote: engineers ran six manual kubectl and migration steps in a wki order, often skipping one. I wrote a single Make target that ran migrations in order, waited on health checks, and posted the deploy SHA to Slack. After rollout, promote time went from about forty-five minutes with frequent rollbacks to twelve minutes, and failed promotes dropped from roughly one in four to near zero over the next month.

Scorable

Yes - None

Dimensions

Follow-up

present=no | probe=none | impact=none | None

Overall Reasoning

None

Flags

None

Case #18: Auto case 18

Rubric: behavioral

Resolved path: behavioral

Overall: 5/5

Case runtime: 25.08s

Question: Tell me about a time you showed leadership when you weren't the formal lead.

Response: During an outage in November 2023 our on-call primary was out sick and I wasn't scheduled but I was online. The incident was elevated payout latency across regions. I joined the bridge, started a shared doc with timeline and hypotheses, and assigned parallel tracks — one person on metrics, one on recent deploys, one on vendor status. I wasn't the manager, but I kept the room focused on disproving theories with data. We found a feature flag rollout that had increased DB connection pool usage; we rolled back within forty minutes of bridge start. Afterward I proposed a pre-flight checklist for flag changes that the team adopted — no title change, but it stuck.

Scorable

Yes - The candidate provides a complete and structured response to a behavioral question.

Dimensions

base (5/5): The response is highly relevant, well-structured, and directly addresses the prompt.

personal_contribution (5/5): The candidate clearly identifies their specific actions, such as starting the doc and assigning tracks, rather than just describing team actions.

real_example (5/5): The candidate provides specific context, including a date, a specific technical issue, and the exact cause found.

outcome (5/5): The candidate describes both an immediate technical resolution and a lasting process improvement.

Follow-up

present=no | probe=none | impact=none | No follow-up probe was provided in the input.

Overall Reasoning

This is an excellent behavioral response that demonstrates leadership through specific, high-impact actions and measurable outcomes.

Flags

None

Case #17: Auto case 17

Rubric: behavioral

Resolved path: behavioral

Overall: 5/5

Case runtime: 26.71s

Question: I hear the framework — can you give me a past example when you actually had to do that?

Response: Yes. In January 2024 we had a production ledger discrepancy, a bank certification deadline, and a hiring loop I was running — all in the same week. I triaged with my manager: the discrepancy was P0 because of money movement, certification was P1 with a fixed external date, recruiting was P2. I spent Monday through Wednesday on the discrepancy root cause — turned out to be a timezone boundary in cutover logic — shipped a patch with verification from finance, then shifted to certification test cases Thursday and Friday. I handed the recruiting debrief to another engineer who had been shadowing the loop. We hit the bank date; the ledger fix held; we still made the hire.

Scorable

Yes - The candidate provides a concrete past example with specific details and outcomes.

Dimensions

base (5/5): The candidate demonstrates excellent prioritization skills by categorizing tasks by urgency and impact.

personal_contribution (5/5): The candidate clearly outlines their specific actions, including root cause analysis and task delegation.

real_example (5/5): The response includes specific dates, technical details, and a clear context.

outcome (5/5): All objectives mentioned in the scenario were successfully achieved.

Follow-up

present=no | probe=none | impact=none | No follow-up probe was provided in the input.

Overall Reasoning

A high-quality behavioral response that covers all necessary dimensions with specificity and clear results.

Flags

None

Case #19: Auto case 19

Rubric: behavioral

Resolved path: behavioral

Overall: 5/5

Case runtime: 27.41s

Question: How do you typically work with product managers — give me a real pattern, not an ideal.

Response: At LedgerPay my PM and I do weekly thirty-minute syncs plus async Slack. I push back early if a spec hand-waves cutoffs or timezone rules — I'll draft the edge-case list before we estimate. When we disagree, I insist we write down assumptions and pick a cheap spike if it's architectural. Last quarter she wanted "real-time reconciliation," which meant different things to us; we spent one sprint on a prototype that proved sub-minute matching wasn't worth the cost for our volume, and we shipped hourly instead with a clear SLA. It's not always smooth, but the habit is: name the tradeoff, measure, then commit.

Scorable

Yes - The candidate provides a specific pattern and a concrete example of working with a PM.

Dimensions

base (5/5): The response is highly structured and directly addresses the prompt's request for a real pattern rather than an ideal.

personal_contribution (5/5): The candidate clearly outlines their specific actions, such as drafting edge-case lists and driving the use of spikes.

real_example (5/5): The reconciliation prototype serves as a strong, specific real-world example of conflict resolution.
outcome (5/5): The example concludes with a clear, data-driven decision that optimized cost and performance via an hourly SLA.

Follow-up

present=no | probe=none | impact=none | None

Overall Reasoning

A top-tier behavioral response that includes a clear pattern, personal agency, a concrete example, and a successful outcome.

Flags

None

Case #20: Auto case 20

Rubric: non_behavioral

Resolved path: skip

Overall: skip

Case runtime: 39.85s

Question: RFC for anything that touches the ledger or event bus; fraud has a standing reviewer on that committee. First ninety days: ship one medium-sized improvement to observability or deploy safety, get comfortable on-call, and build trust with fraud and core banking. On-call we're six engineers per rotation, one week primary and one backup, weekends included but pages have been roughly two per week average this quarter. Anything else?

Response: That's the main list. Maybe one more — is there room to work on the event-sourced ledger itself in this role, or mostly consuming it?

Scorable

No - The candidate is asking a clarifying question about the role's scope rather than providing a response that demonstrates competency.

Dimensions

base (?/5): None

Follow-up

present=no | probe=none | impact=none | None

Overall Reasoning

None

Flags

None

Case #16: Auto case 16

Rubric: behavioral

Resolved path: behavioral

Overall: 1/5

Case runtime: 54.87s

Question: How do you handle competing priorities when everything is urgent?

Response: If I were in that situation, I would list tasks by customer impact and regulatory exposure first, then communicate tradeoffs to product and my manager in writing so there's a shared record. I'd time-box investigation on Sev-2s and escalate when scope creeps.

Scorable

Yes - The candidate provides a direct response to the interviewer's question.

Dimensions

base (3/5): The candidate outlines a logical and professional framework for prioritization and communication.

personal_contribution (1/5): The response is entirely hypothetical and does not describe a specific action taken by the candidate in a real situation.

real_example (1/5): No specific past event or real-world scenario was provided.

outcome (1/5): No result or measurable impact was mentioned as the response was hypothetical.

Follow-up

present=no | probe=none | impact=none | No follow-up was provided in this turn.

Overall Reasoning

The candidate failed to provide a behavioral response, offering only a hypothetical framework instead of demonstrating past experience through a real example.

Flags

hypothetical_response

Case #22: Auto case 22

Rubric: non_behavioral
Resolved path: skip
Overall: skip
Case runtime: 36.02s
Question: Have a good rest of your day.

# Expected rubric mapping (for manual / automated integration checks)			
# Segment	Rubric type	Notes	
# ----- ----- -----			
# Tell me about yourself (short)	non_behavioral	Strong specificity, role-aligned	
# Walk me through your resume (full arc)	non_behavioral	Chronological; concrete roles/metrics	
# Resume probe — PayLoom scaling	non_behavioral	STAR story embedded in resume answer	
# Resume probe — reconciliation vs settlement	non_behavioral	Clarification, technical specificity	
# Gap / overlap question	non_behavioral	Direct, brief	
# Why Northline / this role	non_behavioral	Company-specific, not generic	
# How have you used Go in production	non_behavioral	Skill depth; concrete patterns	
# Tight deadline (settlement migration)	behavioral STAR	Strong; follow-up on replay	
# Persuade skeptical stakeholder (DLQ)	behavioral STAR	Strong	
# Conflict with teammate (testing)	behavioral STAR	Good; minor filler at open	
# Failed / significant mistake (validation)	behavioral STAR	Strong; owns failure + remediation	
# Tough feedback (steering meetings)	behavioral STAR	Strong	
# Improve a process (first answer)	behavioral STAR	Weak — vague until follow-up	
# Improve a process (after follow-up)	behavioral STAR	Stronger	
# Mentoring (first answer)	behavioral STAR	Weak context — missing situation	
# Mentoring (after follow-up)	behavioral STAR	Stronger with LedgerPay context	
# Greatest weakness	non_behavioral	Genuine weakness + remediation	
# Competing priorities (first answer)	behavioral STAR	HYPOTHETICAL — should flag / need follow-up	
# Competing priorities (second answer)	behavioral STAR	Real example after probe	
# Leadership without formal lead (outage)	behavioral STAR	Strong	
# How you work with PMS	non_behavioral	Real pattern; tradeoff example	
# Candidate questions	n/a	Not scored as interview response	

Response: None

Scorable

No - The turn is a closing pleasantry and contains no candidate response to evaluate.

Dimensions

base (?/5): None

Follow-up

present=no | probe=none | impact=none | None

Overall Reasoning

The turn is not scorable.

Flags

None

Case #21: Auto case 21

Rubric: non_behavioral

Resolved path: skip

Overall: skip

Case runtime: 41.18s

Question: Both. Most work is consumption and hardening, but we've had seniors contribute to ordering guarantees and replay tooling. If that's your interest, there's a path. We'll follow up within a week on next steps. Thanks for the thorough walkthrough today, Alex.

Response: Thank you, Priya. I appreciated the depth of the questions — this reinforced that the problems here are the ones I want. Looking forward to hearing from you.

Scorable

No - The turn consists of closing pleasantries and does not contain any substantive content to evaluate.

Dimensions

base (?/5): None

Follow-up

present=no | probe=none | impact=none | None

Overall Reasoning

The turn is a social closing and provides no evaluable signal.

Flags

None